

Flatiron Phase 3 Seasonal Flu Vaccination Project

- Jonah Devoy
- Contact Me: [Linkedin URL \(www.linkedin.com/in/jonahdevoy\)](https://www.linkedin.com/in/jonahdevoy)

Features:

For all binary variables: 0 = No; 1 = Yes.

- h1n1_concern - Level of concern about the H1N1 flu. 0 = Not at all concerned; 1 = Not very concerned; 2 = Somewhat concerned; 3 = Very concerned.
- h1n1_knowledge - Level of knowledge about H1N1 flu. 0 = No knowledge; 1 = A little knowledge; 2 = A lot of knowledge.
- behavioral_antiviral_meds - Has taken antiviral medications. (binary)
- behavioral_avoidance - Has avoided close contact with others with flu-like symptoms. (binary)
- behavioral_face_mask - Has bought a face mask. (binary)
- behavioral_wash_hands - Has frequently washed hands or used hand sanitizer. (binary)
- behavioral_large_gatherings - Has reduced time at large gatherings. (binary)
- behavioral_outside_home - Has reduced contact with people outside of own household. (binary)
- behavioral_touch_face - Has avoided touching eyes, nose, or mouth. (binary)
- doctor_recc_h1n1 - H1N1 flu vaccine was recommended by doctor. (binary)
- doctor_recc_seasonal - Seasonal flu vaccine was recommended by doctor. (binary)
- chronic_med_condition - Has any of the following chronic medical conditions: asthma or an other lung condition, diabetes, a heart condition, a kidney condition, sickle cell anemia or other anemia, a neurological or neuromuscular condition, a liver condition, or a weakened immune system caused by a chronic illness or by medicines taken for a chronic illness. (binary)
- child_under_6_months - Has regular close contact with a child under the age of six months. (binary)
- health_worker - Is a healthcare worker. (binary)
- health_insurance - Has health insurance. (binary)
- opinion_h1n1_vacc_effective - Respondent's opinion about H1N1 vaccine effectiveness. 1 = Not at all effective; 2 = Not very effective; 3 = Don't know; 4 = Somewhat effective; 5 = Very effective.
- opinion_h1n1_risk - Respondent's opinion about risk of getting sick with H1N1 flu without vaccine. 1 = Very Low; 2 = Somewhat low; 3 = Don't know; 4 = Somewhat high; 5 = Very high.
- opinion_h1n1_sick_from_vacc - Respondent's worry of getting sick from taking H1N1 vaccine. 1 = Not at all worried; 2 = Not very worried; 3 = Don't know; 4 = Somewhat worried; 5 = Very worried.
- opinion_seas_vacc_effective - Respondent's opinion about seasonal flu vaccine effectiveness. 1 = Not at all effective; 2 = Not very effective; 3 = Don't know; 4 = Somewhat effective; 5 = Very effective.
- opinion_seas_risk - Respondent's opinion about risk of getting sick with seasonal flu without vaccine. 1 = Very Low; 2 = Somewhat low; 3 = Don't know; 4 = Somewhat high; 5 = Very high.
- opinion_seas_sick_from_vacc - Respondent's worry of getting sick from taking seasonal flu vaccine. 1 = Not at all worried; 2 = Not very worried; 3 = Don't know; 4 = Somewhat worried; 5 = Very worried.
- age_group - Age group of respondent.
- education - Self-reported education level.
- race - Race of respondent.
- sex - Sex of respondent.
- income_poverty - Household annual income of respondent with respect to 2008 Census poverty thresholds.
- marital_status - Marital status of respondent.
- rent_or_own - Housing situation of respondent.
- employment_status - Employment status of respondent.
- hhs_geo_region - Respondent's residence using a 10-region geographic classification defined by the U.S. Dept. of Health and Human Services. Values are represented as short random character strings.
- census_msa - Respondent's residence within metropolitan statistical areas (MSA) as defined by the U.S. Census.
- household_adults - Number of other adults in household, top-coded to 3.
- household_children - Number of children in household, top-coded to 3.
- employment_industry - Type of industry respondent is employed in. Values are represented as short random character strings.
- employment_occupation - Type of occupation of respondent. Values are represented as short random character strings.



```
In [1]: # import necessary libraries

import pandas as pd
pd.set_option('display.max_columns', 100)
pd.set_option('display.float_format', lambda x: '%.3f' % x)

import numpy as np

import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
sns.set(style="darkgrid")

from ydata_profiling import ProfileReport
import missingno
%matplotlib inline
# from IPython.display import Image
## import function needed for split
from sklearn.model_selection import train_test_split

## import classes necessary for building preprocessing pipelines
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler, OneHotEncoder
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.pipeline import Pipeline
from sklearn import tree

from warnings import simplefilter
simplefilter(action='ignore', category=FutureWarning)
```

```
In [2]: def calculate_null_percentage(df):

    import pandas as pd

    missing_vals = pd.DataFrame()
    missing_vals['Number of Nulls'] = df.isna().sum()
    missing_vals['% Null'] = (df.isna().sum() / len(df)) * 100

    return missing_vals

def check_unique(df, col, dropna=False):

    import pandas as pd

    unique_vals = pd.DataFrame(df[col].value_counts(dropna=dropna))

    return unique_vals
```

```
In [3]: #conda install -c conda-forge missingno
```

```
In [4]: #pip install -U ydata-profiling
```

```
In [5]: # Loading in the data
df_1 = pd.read_csv('/Users/jdapeman/Documents/flu_shot_V1/CSV_FOLDER/Flu_Shot_Learning_Predict_H1N1_and_Seasonal_Flu_Va
features_df = pd.read_csv('/Users/jdapeman/Documents/flu_shot_V1/CSV_FOLDER/Flu_Shot_Learning_Predict_H1N1_and_Seasonal
labels_df = pd.read_csv('/Users/jdapeman/Documents/flu_shot_V1/CSV_FOLDER/Flu_Shot_Learning_Predict_H1N1_and_Seasonal_F
```

```
In [6]: #profile1 = ProfileReport(df, title="Profiling Report")
#profile1
```

```
In [7]: #profile2 = ProfileReport(features_df, title="Profiling Report")
#profile2
```



```
In [8]: #profile3 = ProfileReport(labels_df, title="Profiling Report")
#profile3
```

```
In [9]: # observing the first 5 rows of the features dataframe
features_df.head()
```

```
Out[9]:
```

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavioral_large
0	0	1.000	0.000	0.000	0.000	0.000	0.000	
1	1	3.000	2.000	0.000	1.000	0.000	1.000	
2	2	1.000	1.000	0.000	1.000	0.000	0.000	
3	3	1.000	1.000	0.000	1.000	0.000	1.000	
4	4	2.000	1.000	0.000	1.000	0.000	1.000	

```
In [10]: # observing the first 5 rows of the labels dataframe
labels_df.head()
```

```
Out[10]:
```

	respondent_id	h1n1_vaccine	seasonal_vaccine
0	0	0	0
1	1	0	1
2	2	0	0
3	3	0	1
4	4	0	0

```
In [11]: # observing the first 5 rows of the test features dataframe
df_1.head()
```

```
Out[11]:
```

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavioral_large
0	26707	2.000	2.000	0.000	1.000	0.000	1.000	
1	26708	1.000	1.000	0.000	0.000	0.000	0.000	
2	26709	2.000	2.000	0.000	0.000	1.000	1.000	
3	26710	1.000	1.000	0.000	0.000	0.000	0.000	
4	26711	3.000	1.000	1.000	1.000	0.000	1.000	



```
In [12]: # combining the tables and training features into a single dataframe
df = pd.concat([features_df, labels_df.drop('respondent_id', axis=1)], axis=1)
df
```

Out[12]:

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavioral_
0	0	1.000	0.000	0.000	0.000	0.000	0.000	
1	1	3.000	2.000	0.000	1.000	0.000	1.000	
2	2	1.000	1.000	0.000	1.000	0.000	0.000	
3	3	1.000	1.000	0.000	1.000	0.000	1.000	
4	4	2.000	1.000	0.000	1.000	0.000	1.000	
...	
26702	26702	2.000	0.000	0.000	1.000	0.000	0.000	
26703	26703	1.000	2.000	0.000	1.000	0.000	1.000	
26704	26704	2.000	2.000	0.000	1.000	1.000	1.000	
26705	26705	1.000	1.000	0.000	0.000	0.000	0.000	
26706	26706	0.000	0.000	0.000	1.000	0.000	0.000	

26707 rows x 38 columns

```
In [13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26707 entries, 0 to 26706
Data columns (total 38 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   respondent_id                             26707 non-null  int64
1   h1n1_concern                             26615 non-null  float64
2   h1n1_knowledge                           26591 non-null  float64
3   behavioral_antiviral_meds                 26636 non-null  float64
4   behavioral_avoidance                     26499 non-null  float64
5   behavioral_face_mask                     26688 non-null  float64
6   behavioral_wash_hands                     26665 non-null  float64
7   behavioral_large_gatherings               26620 non-null  float64
8   behavioral_outside_home                   26625 non-null  float64
9   behavioral_touch_face                     26579 non-null  float64
10  doctor_recc_h1n1                         24547 non-null  float64
11  doctor_recc_seasonal                     24547 non-null  float64
12  chronic_med_condition                     25736 non-null  float64
13  child_under_6_months                     25887 non-null  float64
14  health_worker                             25903 non-null  float64
15  health_insurance                         14433 non-null  float64
16  opinion_h1n1_vacc_effective                26316 non-null  float64
17  opinion_h1n1_risk                          26319 non-null  float64
18  opinion_h1n1_sick_from_vacc                26312 non-null  float64
19  opinion_seas_vacc_effective                 26245 non-null  float64
20  opinion_seas_risk                          26193 non-null  float64
21  opinion_seas_sick_from_vacc                 26170 non-null  float64
22  age_group                                 26707 non-null  object
23  education                                 25300 non-null  object
24  race                                      26707 non-null  object
25  sex                                       26707 non-null  object
26  income_poverty                           22284 non-null  object
27  marital_status                           25299 non-null  object
28  rent_or_own                              24665 non-null  object
29  employment_status                        25244 non-null  object
30  hhs_geo_region                           26707 non-null  object
31  census_msa                               26707 non-null  object
32  household_adults                         26458 non-null  float64
33  household_children                       26458 non-null  float64
34  employment_industry                      13377 non-null  object
35  employment_occupation                    13237 non-null  object
36  h1n1_vaccine                             26707 non-null  int64
37  seasonal_vaccine                         26707 non-null  int64
dtypes: float64(23), int64(3), object(12)
memory usage: 7.7+ MB
```



```
In [14]: df.describe()
```

```
Out[14]:
```

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavioral_w
count	26707.000	26615.000	26591.000	26636.000	26499.000	26688.000	26665.000	
mean	13353.000	1.618	1.263	0.049	0.726	0.069	0.826	
std	7709.791	0.910	0.618	0.216	0.446	0.253	0.379	
min	0.000	0.000	0.000	0.000	0.000	0.000	0.000	
25%	6676.500	1.000	1.000	0.000	0.000	0.000	1.000	
50%	13353.000	2.000	1.000	0.000	1.000	0.000	1.000	
75%	20029.500	2.000	2.000	0.000	1.000	0.000	1.000	
max	26706.000	3.000	2.000	1.000	1.000	1.000	1.000	

```
In [15]: df.corr(numeric_only = [False])
```

```
Out[15]:
```

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_w
respondent_id	1.000	0.018	0.003	-0.008	0.010	-0.007	
h1n1_concern	0.018	1.000	0.063	0.090	0.234	0.156	
h1n1_knowledge	0.003	0.063	1.000	-0.011	0.089	0.030	
behavioral_antiviral_meds	-0.008	0.090	-0.011	1.000	0.049	0.146	
behavioral_avoidance	0.010	0.234	0.089	0.049	1.000	0.065	
behavioral_face_mask	-0.007	0.156	0.030	0.146	0.065	1.000	
behavioral_wash_hands	0.011	0.294	0.090	0.064	0.338	0.083	
behavioral_large_gatherings	0.005	0.255	-0.049	0.106	0.228	0.181	
behavioral_outside_home	0.009	0.246	-0.068	0.128	0.220	0.163	
behavioral_touch_face	0.008	0.248	0.086	0.071	0.335	0.104	
doctor_recc_h1n1	-0.002	0.150	0.094	0.051	0.068	0.084	
doctor_recc_seasonal	0.001	0.136	0.072	0.031	0.074	0.069	
chronic_med_condition	0.006	0.095	-0.023	0.008	0.039	0.068	
child_under_6_months	-0.005	0.050	0.022	0.029	-0.000	0.040	
health_worker	-0.003	0.034	0.170	0.009	0.001	0.070	
health_insurance	-0.013	-0.004	0.119	-0.064	0.033	-0.040	
opinion_h1n1_vacc_effective	0.006	0.240	0.121	0.030	0.112	0.038	
opinion_h1n1_risk	0.001	0.377	0.073	0.105	0.118	0.131	
opinion_h1n1_sick_from_vacc	-0.002	0.360	-0.020	0.079	0.131	0.107	
opinion_seas_vacc_effective	0.006	0.235	0.086	0.015	0.120	0.042	
opinion_seas_risk	-0.005	0.334	0.077	0.085	0.130	0.110	
opinion_seas_sick_from_vacc	0.010	0.226	-0.062	0.084	0.083	0.090	
household_adults	0.000	-0.016	0.025	0.045	0.019	0.014	
household_children	-0.004	0.051	0.051	0.085	0.040	0.006	
h1n1_vaccine	-0.003	0.122	0.118	0.041	0.048	0.070	
seasonal_vaccine	-0.005	0.155	0.120	0.006	0.076	0.050	

```
In [16]: df.shape
```

```
Out[16]: (26707, 38)
```



```
In [17]: # Check for null or missing values in each column
for column in df.columns:
    null_values = df[column].isnull()

    if null_values.any():
        print(f"Null values in the column '{column}':")
        print(df[column][null_values])
        print("=====")
```

Null values in the column 'h1n1_concern':

```
44      NaN
96      NaN
150     NaN
411     NaN
758     NaN
..
25788   NaN
25883   NaN
25948   NaN
26358   NaN
26471   NaN
```

Name: h1n1_concern, Length: 92, dtype: float64

=====

Null values in the column 'h1n1_knowledge':

```
136      NaN
405      NaN
958      NaN
1026     NaN
....
.. ..
```

```
In [18]: # the focus for this notebook is on seasonal flu vaccines, so any columns relating to H1N1 vaccines can be dropped
# and we will drop columns unrelated to the seasonal flu vaccine
df.drop(columns=['opinion_h1n1_vacc_effective',
                 'opinion_h1n1_risk',
                 'opinion_h1n1_sick_from_vacc',
                 'doctor_recc_h1n1',
                 'h1n1_vaccine',
                 'behavioral_antiviral_meds',
                 'respondent_id'], axis=1, inplace=True)
```

```
In [19]: df.shape
```

```
Out[19]: (26707, 31)
```

```
In [20]: df.head()
```

```
Out[20]:
```

	h1n1_concern	h1n1_knowledge	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavioral_large_gatherings	behavioral_outside_home	be
0	1.000	0.000	0.000	0.000	0.000	0.000	1.000	
1	3.000	2.000	1.000	0.000	1.000	0.000	1.000	
2	1.000	1.000	1.000	0.000	0.000	0.000	0.000	
3	1.000	1.000	1.000	0.000	1.000	1.000	0.000	
4	2.000	1.000	1.000	0.000	1.000	1.000	0.000	



In [21]: `df.info()`

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26707 entries, 0 to 26706
Data columns (total 31 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   h1n1_concern                             26615 non-null  float64
1   h1n1_knowledge                           26591 non-null  float64
2   behavioral_avoidance                     26499 non-null  float64
3   behavioral_face_mask                     26688 non-null  float64
4   behavioral_wash_hands                     26665 non-null  float64
5   behavioral_large_gatherings              26620 non-null  float64
6   behavioral_outside_home                  26625 non-null  float64
7   behavioral_touch_face                     26579 non-null  float64
8   doctor_recc_seasonal                     24547 non-null  float64
9   chronic_med_condition                   25736 non-null  float64
10  child_under_6_months                     25887 non-null  float64
11  health_worker                           25903 non-null  float64
12  health_insurance                         14433 non-null  float64
13  opinion_seas_vacc_effective               26245 non-null  float64
14  opinion_seas_risk                         26193 non-null  float64
15  opinion_seas_sick_from_vacc               26170 non-null  float64
16  age_group                               26707 non-null  object
17  education                               25300 non-null  object
18  race                                    26707 non-null  object
19  sex                                    26707 non-null  object
20  income_poverty                          22284 non-null  object
21  marital_status                          25299 non-null  object
22  rent_or_own                             24665 non-null  object
23  employment_status                       25244 non-null  object
24  hhs_geo_region                          26707 non-null  object
25  census_msa                              26707 non-null  object
26  household_adults                        26458 non-null  float64
27  household_children                      26458 non-null  float64
28  employment_industry                     13377 non-null  object
29  employment_occupation                   13237 non-null  object
30  seasonal_vaccine                        26707 non-null  int64
dtypes: float64(18), int64(1), object(12)
memory usage: 6.3+ MB

```

The columns `hhs_geo_region`, `employment_industry`, and `employment_occupation` are random strings that have been scrambled and encoded for anonymity, the definition is not provided by the CDC. But since they are consistent and specific to the individual, they can be used to make accurate predictions.

In [22]: `#pip install pandasgui`In [23]: `#pip install git+https://github.com/adameroose/pandasgui.git`In [24]: `# pandas gui
#from pandasgui import show
#show(df)`In [25]: `# A function to print out the number of nulls in each column as well as the percentage of nulls
def calculate_null_percentage(df):
 missing_vals = pd.DataFrame()
 missing_vals['Number of Nulls'] = df.isna().sum()
 missing_vals['% Null'] = (df.isna().sum() / len(df)) * 100

 return missing_vals`

```
In [26]: calculate_null_percentage(df)
```

```
Out[26]:
```

	Number of Nulls	% Null
h1n1_concern	92	0.344
h1n1_knowledge	116	0.434
behavioral_avoidance	208	0.779
behavioral_face_mask	19	0.071
behavioral_wash_hands	42	0.157
behavioral_large_gatherings	87	0.326
behavioral_outside_home	82	0.307
behavioral_touch_face	128	0.479
doctor_recc_seasonal	2160	8.088
chronic_med_condition	971	3.636
child_under_6_months	820	3.070
health_worker	804	3.010
health_insurance	12274	45.958
opinion_seas_vacc_effective	462	1.730
opinion_seas_risk	514	1.925
opinion_seas_sick_from_vacc	537	2.011
age_group	0	0.000
education	1407	5.268
race	0	0.000
sex	0	0.000
income_poverty	4423	16.561
marital_status	1408	5.272
rent_or_own	2042	7.646
employment_status	1463	5.478
hhs_geo_region	0	0.000
census_msa	0	0.000
household_adults	249	0.932
household_children	249	0.932
employment_industry	13330	49.912
employment_occupation	13470	50.436
seasonal_vaccine	0	0.000

```
In [27]: # Calculate the percentage of missing values for each column and then
# filter columns with less than 5% missing values and not equal to zero
def print_columns_missing_info(df):
    total_rows = len(df)
    missing_percentages = (df.isnull().sum() / total_rows) * 100
    columns_missing_info = missing_percentages[(missing_percentages < 5) & (missing_percentages > 0)]
    if len(columns_missing_info) > 0:
        for column in columns_missing_info.index:
            print(column)

print_columns_missing_info(df)
```

```
h1n1_concern
h1n1_knowledge
behavioral_avoidance
behavioral_face_mask
behavioral_wash_hands
behavioral_large_gatherings
behavioral_outside_home
behavioral_touch_face
chronic_med_condition
child_under_6_months
health_worker
opinion_seas_vacc_effective
opinion_seas_risk
opinion_seas_sick_from_vacc
household_adults
household_children
```



Notes About the Data and going forward:

- We can see the columns employment_occupation, employment_industry, and health_insurance exhibit the highest number of missing values. Among the null values in employment_occupation and employment_industry, 10,231 values correspond to individuals categorized as 'Not in Labor Force' in the employment_status column. These can be considered as N/A rather than individuals declining to answer. 1,453 observances represent unemployed individuals where the employment_occupation and employment_industry columns are appropriately labeled as 'not applicable'.
- We can also see that if an individual declined to answer if their doctor recommended one type of vaccine, they typically refused to answer about the recommendation for the other type. There is a tendency for individuals to refuse to answer questions related to having a chronic medical condition, having a child under 6 months, being a health worker, having opinion-based questions, income, education, and personal and home life topics, for reasons unknown to us.
- We should be treating missing information for specific variables as its distinctive category, instead of removing it because a non answer in itself is a kind of response

```
In [28]: no_null_cols = [col for col in df.columns if df[col].isna().sum()==0]
no_null_cols
```

```
Out[28]: ['age_group',
'race',
'sex',
'hhs_geo_region',
'census_msa',
'seasonal_vaccine']
```

```
In [29]: # select individuals not in the labor force
not_in_labor_force = df[df['employment_status']=='Not in Labor Force']
calculate_null_percentage(not_in_labor_force)
```

	Number of Nulls	% Null
h1n1_concern	56	0.547
h1n1_knowledge	58	0.567
behavioral_avoidance	104	1.017
behavioral_face_mask	9	0.088
behavioral_wash_hands	22	0.215
behavioral_large_gatherings	46	0.450
behavioral_outside_home	49	0.479
behavioral_touch_face	67	0.655
doctor_recc_seasonal	843	8.240
chronic_med_condition	91	0.889
child_under_6_months	1	0.010
health_worker	6	0.059

```
In [30]: # select unemployed individuals
unemployed = df[df['employment_status']=='Unemployed']
calculate_null_percentage(unemployed)
```

```
Out[30]:
```

	Number of Nulls	% Null
h1n1_concern	3	0.206
h1n1_knowledge	10	0.688
behavioral_avoidance	6	0.413
behavioral_face_mask	1	0.069
behavioral_wash_hands	0	0.000
behavioral_large_gatherings	5	0.344
behavioral_outside_home	2	0.138
behavioral_touch_face	7	0.482
doctor_recc_seasonal	94	6.469
chronic_med_condition	16	1.101
child_under_6_months	1	0.069

```
In [31]: # takes the dataframe and column name and returns the unique values in that column as well as the
# number of each unique values.
def count_unique_values(df, col, dropna=False):
    unique_vals = pd.DataFrame(df[col].value_counts(dropna=dropna))

    return unique_vals
```



```
In [32]: # creating not_employed from labor force
# if a person is unemployed change their 'employment_industry' to 'not_employed'
df.loc[df['employment_status'] == 'Unemployed', 'employment_industry'] = 'not_employed'

# if a person is not in the labor force change their 'employment_industry' to 'not_employed'
df.loc[df['employment_status'] == 'Not in Labor Force', 'employment_industry'] = 'not_employed'

count_unique_values(df, 'employment_industry')
```

Out[32]:

employment_industry	
not_employed	11684
fcxhlnwr	2468
wxleyezf	1804
NaN	1646
ldnlellj	1231
pxcmvdjn	1037
atmlpfrs	926
arjwrjbj	871
xicduogh	851
mfikgejo	614
vjjrobsf	527
rucpzij	523
xqicxuve	511
saaquncn	338
cfqqtusy	325
nduyfdeo	286
mcubkhph	275
wlfvacwt	215
dotnnunm	201
haxffmxo	148
msuufmds	124
phxvnwax	89
qnlwzans	13



```
In [33]: # creating not_employed from employment industry
# if a person is unemployed, change their 'employment_industry' to 'not_employed'
df.loc[df['employment_status'] == 'Unemployed', 'employment_occupation'] = 'not employed'

# if a person is not in the labor force, change their 'employment_industry' to 'not_employed'
df.loc[df['employment_status'] == 'Not in Labor Force', 'employment_occupation'] = 'not employed'

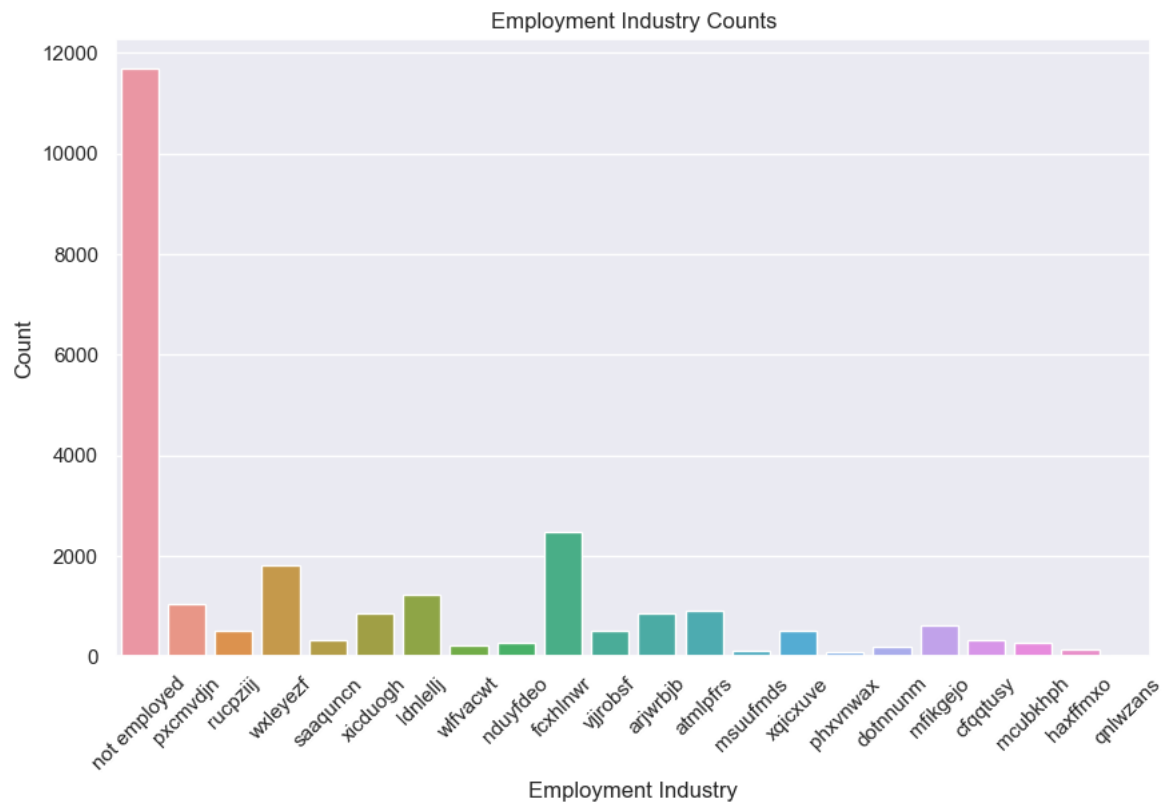
count_unique_values(df, 'employment_occupation')
```

Out[33]:

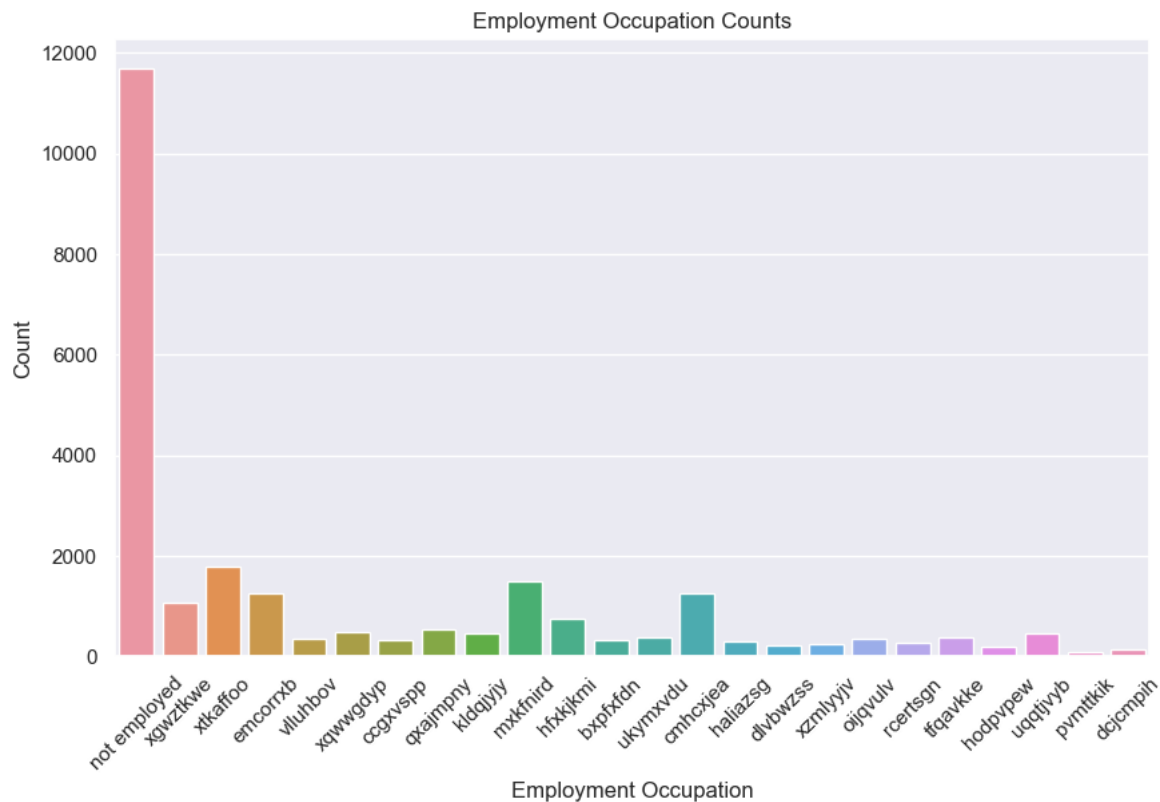
employment_occupation	
not employed	11684
NaN	1786
xtkaffoo	1778
mxkfnird	1509
emcorrx	1270
cmhcxjea	1247
xgwztkwe	1082
hfxkjkmi	766
qxajmpny	548
xqwwgdyp	485
kldqjyiy	469
uqqtjvyb	452
tfqavkke	388
ukymxvdu	372
vluhbov	354
oijqvulv	344
cogxvspp	341
bxpfxfdn	331
haliazsg	296
rcertsgn	276
xzmlyyiv	248
dlvbwzss	227
hodpvpew	208
dcjcpih	148
pvmttkik	98



```
In [34]: # Bar plot of employment industry counts
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='employment_industry')
plt.title('Employment Industry Counts')
plt.xlabel('Employment Industry')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.show()
```



```
In [35]: # Bar plot of employment occupation counts
plt.figure(figsize=(10, 6))
sns.countplot(data=df, x='employment_occupation')
plt.title('Employment Occupation Counts')
plt.xlabel('Employment Occupation')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.show()
```



```
In [36]: # columns with null between 0 and 5%, excluding values of 0
null_df = calculate_null_percentage(df)
null_df.drop(index=null_df.loc[null_df['% Null']==0].index, axis=0, inplace=True)
under_5_null = null_df.loc[null_df['% Null']<5]
under_5_null
```

Out[36]:

	Number of Nulls	% Null
h1n1_concern	92	0.344
h1n1_knowledge	116	0.434
behavioral_avoidance	208	0.779
behavioral_face_mask	19	0.071
behavioral_wash_hands	42	0.157
behavioral_large_gatherings	87	0.326
behavioral_outside_home	82	0.307
behavioral_touch_face	128	0.479
chronic_med_condition	971	3.636
child_under_6_months	820	3.070
health_worker	804	3.010
opinion_seas_vacc_effective	462	1.730
opinion_seas_risk	514	1.925
opinion_seas_sick_from_vacc	537	2.011
household_adults	249	0.932
household_children	249	0.932



```
In [37]: # dropping the rows with less than 5% of null values
under_5_null_cols = list(under_5_null.index)
df.dropna(subset=under_5_null_cols, inplace=True)
df.head()
```

```
Out[37]:
```

	h1n1_concern	h1n1_knowledge	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavioral_large_gatherings	behavioral_outside_home	be
0	1.000	0.000	0.000	0.000	0.000	0.000	1.000	
1	3.000	2.000	1.000	0.000	1.000	0.000	1.000	
2	1.000	1.000	1.000	0.000	0.000	0.000	0.000	
3	1.000	1.000	1.000	0.000	1.000	1.000	0.000	
4	2.000	1.000	1.000	0.000	1.000	1.000	0.000	

```
In [38]: df.shape
```

```
Out[38]: (24939, 31)
```


```
In [39]: # the result is a dataframe with rows containing less than 5% of info have been dropped
calculate_null_percentage(df)
```

```
Out[39]:
```


	Number of Nulls	% Null
h1n1_concern	0	0.000
h1n1_knowledge	0	0.000
behavioral_avoidance	0	0.000
behavioral_face_mask	0	0.000
behavioral_wash_hands	0	0.000
behavioral_large_gatherings	0	0.000
behavioral_outside_home	0	0.000
behavioral_touch_face	0	0.000
doctor_recc_seasonal	1851	7.422
chronic_med_condition	0	0.000
child_under_6_months	0	0.000
health_worker	0	0.000
health_insurance	11043	44.280
opinion_seas_vacc_effective	0	0.000
opinion_seas_risk	0	0.000
opinion_seas_sick_from_vacc	0	0.000
age_group	0	0.000
education	530	2.125
race	0	0.000
sex	0	0.000
income_poverty	3318	13.304
marital_status	530	2.125
rent_or_own	1123	4.503
employment_status	578	2.318
hhs_geo_region	0	0.000
census_msa	0	0.000
household_adults	0	0.000
household_children	0	0.000
employment_industry	753	3.019
employment_occupation	879	3.525
seasonal_vaccine	0	0.000



```
In [40]: missingno.bar(df)
```



```
In [41]: cols_without_null = df.columns[df.isnull().sum() == 0].tolist()
missingno.bar(df.drop(columns=cols_without_null))
```



```
In [42]: ## create a list of cols without any null values to be dropped from missingno.matrix
no_null_cols = [col for col in df.columns if df[col].isna().sum() == 0]

no_null_cols
```

```
Out[42]: ['h1n1_concern',
'h1n1_knowledge',
'behavioral_avoidance',
'behavioral_face_mask',
'behavioral_wash_hands',
'behavioral_large_gatherings',
'behavioral_outside_home',
'behavioral_touch_face',
'chronic_med_condition',
'child_under_6_months',
'health_worker',
'opinion_seas_vacc_effective',
'opinion_seas_risk',
'opinion_seas_sick_from_vacc',
'age_group',
'race',
'sex',
'hhs_geo_region',
'census_msa',
'household_adults',
'household_children',
'seasonal_vaccine']
```

- Because 44.2% declined to tell if they have health_insurance, information will become its own category for this variable. The same for the income_poverty variable, which is missing 13.4% of its values.
- The pattern of null values is so similar across the variables education, marital_status, rent_to_own, employment_status, employment_industry, and employment_occupation that dropping rows containing null values for any of these columns will drop most of the records containing null values for the other columns. Null values for these categorical variables will be filled with 'missing' as its own category since this missing information appears to represent a distinct kind of survey respondent.
- The race data is mostly white, with 19,856 out of 24,939 people sampled. The three other categorical races are Black, Hispanic, and Other or Multiple. We will combine these latter three underrepresented groups into one group for people of color.
- The sample population is slightly skewed towards women (59.6%).
- About half of the individuals surveyed declined to answer whether or not they had health insurance.

```
In [43]: #profile_new_V1 = ProfileReport(df, title="Profiling Report")
#profile_new_V1
```

- We are going to make a variable 'behav_score' that represents how much an individual has done to avoid the flu by summing up all behavioral variables. Taking the sum across these columns, a higher score represents a more cautious person.
- We are going to make a variable 'behav_to_risk' that describes the ratio of how much a person has done to avoid the flu in considering their perception of the risk of getting the flu without the vaccine. The numerator is behav_score + 1 to distinguish people who are not doing anything to avoid the flu but vary by how concerned they are about getting sick without the vaccine. The denominator is the rating of risk perception, opinion_seas_risk. An individual with a very low score is someone who has done little to avoid the flu but is very concerned about getting sick without the vaccine. An individual with a score on the upper end has done a lot to behaviorally minimize their risk of exposure and is not very concerned about getting sick without the vaccine. This kind of person may be less likely to get the vaccine, even if they think it's effective because they feel they're doing enough to avoid getting the flu on their own

```
In [44]: behavior_cols = [x for x in df.columns if 'behavioral' in x]
behavior_cols
```

```
Out[44]: ['behavioral_avoidance',
'behavioral_face_mask',
'behavioral_wash_hands',
'behavioral_large_gatherings',
'behavioral_outside_home',
'behavioral_touch_face']
```

```
In [45]: df['behav_score'] = df[behavior_cols].sum(axis=1)
df.columns
```

```
Out[45]: Index(['h1n1_concern', 'h1n1_knowledge', 'behavioral_avoidance',
'behavioral_face_mask', 'behavioral_wash_hands',
'behavioral_large_gatherings', 'behavioral_outside_home',
'behavioral_touch_face', 'doctor_recc_seasonal',
'chronic_med_condition', 'child_under_6_months', 'health_worker',
'health_insurance', 'opinion_seas_vacc_effective', 'opinion_seas_risk',
'opinion_seas_sick_from_vacc', 'age_group', 'education', 'race', 'sex',
'income_poverty', 'marital_status', 'rent_or_own', 'employment_status',
'hhs_geo_region', 'census_msa', 'household_adults',
'household_children', 'employment_industry', 'employment_occupation',
'seasonal_vaccine', 'behav_score'],
dtype='object')
```




```

In [46]: # Get unique values and their counts
def check_column(df, column_name):
    unique_values = df[column_name].value_counts().reset_index()
    unique_values.columns = [column_name, 'Count']

    fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(12, 6))

    sns.histplot(df[column_name].dropna(), kde=False, ax=axes[0])
    axes[0].set_xlabel(column_name)
    axes[0].set_ylabel('Frequency')
    axes[0].set_title('Distribution of {}'.format(column_name))

    sns.boxplot(x=column_name, data=df, ax=axes[1])
    axes[1].set_xlabel(column_name)
    axes[1].set_ylabel('Value')
    axes[1].set_title('Box and Whisker Plot of {}'.format(column_name))

    plt.tight_layout()
    plt.show()


    return unique_values

```

```

In [47]: check_column(df, 'behav_score')

```



```

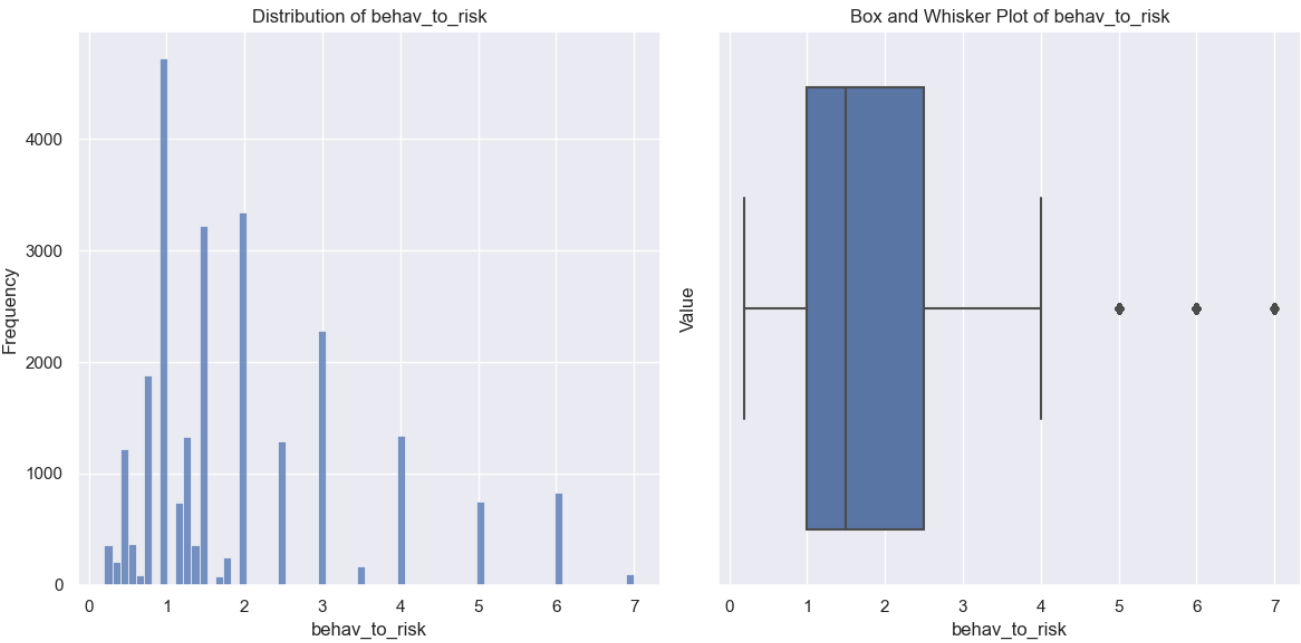
Out[47]:

```

	behav_score	Count
0	3.000	6952
1	5.000	4440
2	2.000	4348
3	4.000	4052
4	1.000	2465
5	0.000	1935
6	6.000	747



```
In [48]: df['behav_to_risk'] = (df['behav_score'] + 1) / df['opinion_seas_risk']
check_column(df, 'behav_to_risk')
```

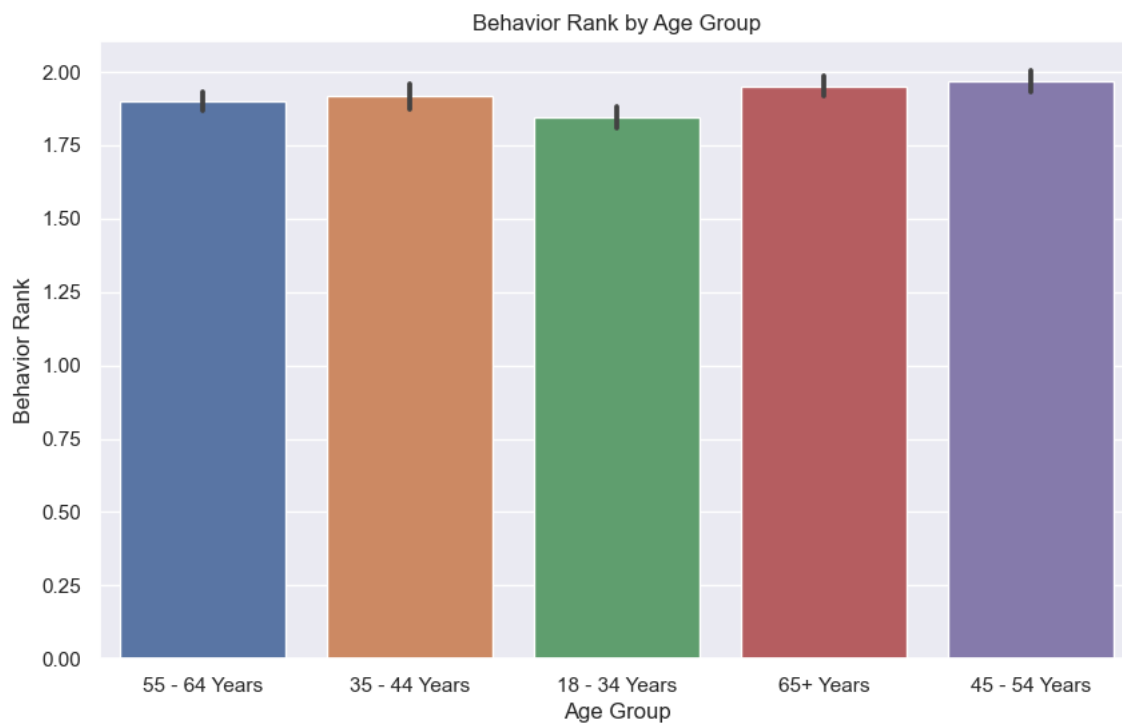


Out[48]:

	behav_to_risk	Count
0	1.000	4721
1	2.000	3345
2	1.500	3222
3	3.000	2281
4	4.000	1344
5	1.250	1326
6	2.500	1290
7	0.500	1218
8	0.750	1179
9	6.000	829
10	5.000	748
11	1.200	742
12	0.800	700
13	0.600	365
14	0.250	283
15	1.750	254
16	1.400	217
17	3.500	168
18	0.400	146
19	1.333	142
20	7.000	98
21	0.667	91
22	1.667	77
23	0.200	76
24	0.333	67
25	2.333	10



```
In [49]: # Plotting 'behavior_rank' and 'age_group'
plt.figure(figsize=(10, 6))
sns.barplot(x='age_group', y='behav_to_risk', data=df)
plt.xlabel('Age Group')
plt.ylabel('Behavior Rank')
plt.title('Behavior Rank by Age Group')
plt.show()
```




```
In [50]: # Create a function for whether or not an individual is 65 years or older as this
# represents a group at higher risk for serious complications from the flu.
def is_older_65(row):
    if row['age_group'] == '65+ Years':
        return 1
    else:
        return 0
```



```
In [51]: # 'older_65' variable for whether or not an individual is 65 years or older as this represents a group at higher risk
# for serious complications from the flu.
df['older_65'] = df.apply(lambda x: is_older_65(x), axis=1)

# check counts of unique values in new col and plot distribution
check_column(df, 'older_65')
```



Out[51]:

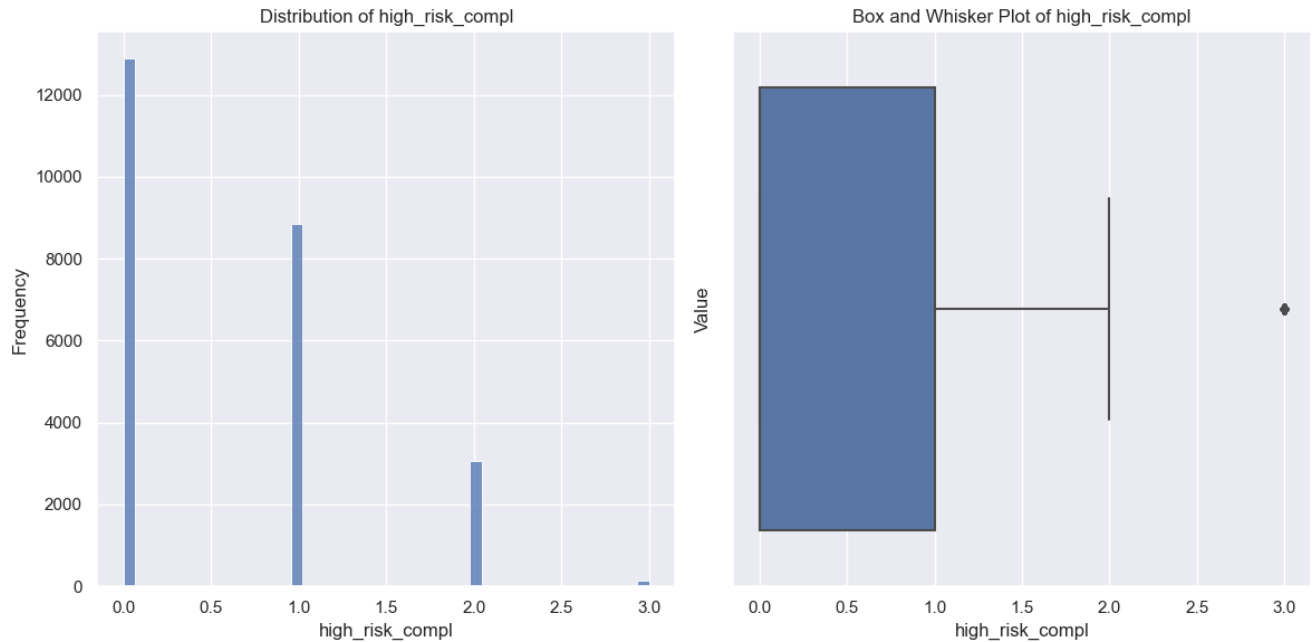
	older_65	Count
0	0	18678
1	1	6261

- Create a variable 'high_risk_compl' if an individual's overall risk for developing flu-related complications. According to the CDC, people 65 years and older, children 6 months or younger, and people with chronic medical conditions are at higher risk for the flu

```
In [52]: # function to calculate score for high risk of complications
def calc_high_risk(row):
    risk = 0
    if row['older_65'] == 1:
        risk += 1
    if row['child_under_6_months'] == 1:
        risk += 1
    if row['chronic_med_condition'] == 1:
        risk += 1
    return risk
```



```
In [53]: # create new column 'high_risk_compl'
df['high_risk_compl'] = df.apply(lambda x: calc_high_risk(x), axis=1)
check_column(df, 'high_risk_compl')
```



```
Out[53]:
```

high_risk_compl	Count
0	12894
1	8852
2	3051
3	142

Making a categorical variable that bins people with multiple high-risk factors (`high_risk_compl > 1`) into one 'high risk' category, assigning 0 to 'low risk' and 1 to 'med risk' Because of the high variations of risk factors per person

```
In [54]: df['high_risk_cat'] = df['high_risk_compl'].map({0:'low risk', 1:'med risk', 2:'high risk', 3:'high risk'})
df['high_risk_cat'].value_counts()
```

```
Out[54]: low risk      12894
med risk      8852
high risk     3193
Name: high_risk_cat, dtype: int64
```

```
In [55]: df['doctor_recc_seasonal'] = df['doctor_recc_seasonal'].map({1.0: '1', 0.0: '0'})
df['doctor_recc_seasonal'].value_counts(dropna=False)
```

```
Out[55]: 0      15420
1       7668
NaN      1851
Name: doctor_recc_seasonal, dtype: int64
```

```
In [56]: df['health_insurance'] = df['health_insurance'].map({1.0: '1', 0.0: '0'})
df['health_insurance'].value_counts(dropna=False)
```

```
Out[56]: 1      12224
NaN     11043
0       1672
Name: health_insurance, dtype: int64
```

```
In [57]: # define a function to return make combine people of color
def race_func(row):
    if row['race'] == 'White':
        return 'White'
    else:
        return 'POC'
```



```
In [58]: df['race'] = df.apply(lambda x: race_func(x), axis=1)
df['race'].value_counts(dropna=False)
```

```
Out[58]: White      19856
POC              5083
Name: race, dtype: int64
```

```
In [59]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 24939 entries, 0 to 26706
Data columns (total 36 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   h1n1_concern                          24939 non-null  float64
1   h1n1_knowledge                        24939 non-null  float64
2   behavioral_avoidance                  24939 non-null  float64
3   behavioral_face_mask                  24939 non-null  float64
4   behavioral_wash_hands                  24939 non-null  float64
5   behavioral_large_gatherings            24939 non-null  float64
6   behavioral_outside_home                24939 non-null  float64
7   behavioral_touch_face                  24939 non-null  float64
8   doctor_recc_seasonal                  23088 non-null  object
9   chronic_med_condition                 24939 non-null  float64
10  child_under_6_months                  24939 non-null  float64
11  health_worker                         24939 non-null  float64
12  health_insurance                      13896 non-null  object
13  opinion_seas_vacc_effective             24939 non-null  float64
14  opinion_seas_risk                       24939 non-null  float64
15  opinion_seas_sick_from_vacc             24939 non-null  float64
16  age_group                             24939 non-null  object
17  education                             24409 non-null  object
18  race                                  24939 non-null  object
19  sex                                   24939 non-null  object
20  income_poverty                        21621 non-null  object
21  marital_status                        24409 non-null  object
22  rent_or_own                           23816 non-null  object
23  employment_status                     24361 non-null  object
24  hhs_geo_region                        24939 non-null  object
25  census_msa                            24939 non-null  object
26  household_adults                      24939 non-null  float64
27  household_children                    24939 non-null  float64
28  employment_industry                   24186 non-null  object
29  employment_occupation                 24060 non-null  object
30  seasonal_vaccine                      24939 non-null  int64
31  behav_score                           24939 non-null  float64
32  behav_to_risk                         24939 non-null  float64
33  older_65                             24939 non-null  int64
34  high_risk_compl                       24939 non-null  int64
35  high_risk_cat                         24939 non-null  object
dtypes: float64(18), int64(3), object(15)
memory usage: 7.0+ MB
```

```
In [60]: # create df with remaining null values filled in with 'missing' for vizualizations
df_missing = df.fillna(value='missing')
df_missing.head()
```

```
Out[60]:
```

	h1n1_concern	h1n1_knowledge	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavioral_large_gatherings	behavioral_outside_home	be
0	1.000	0.000	0.000	0.000	0.000	0.000	1.000	
1	3.000	2.000	1.000	0.000	1.000	0.000	1.000	
2	1.000	1.000	1.000	0.000	0.000	0.000	0.000	
3	1.000	1.000	1.000	0.000	1.000	1.000	0.000	
4	2.000	1.000	1.000	0.000	1.000	1.000	0.000	



```
In [61]: df.corr()
```

```
Out[61]:
```

	h1n1_concern	h1n1_knowledge	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavioral_large_gatherings	b
h1n1_concern	1.000	0.060	0.236	0.159	0.297	0.258	
h1n1_knowledge	0.060	1.000	0.082	0.035	0.087	-0.047	
behavioral_avoidance	0.236	0.082	1.000	0.063	0.337	0.231	
behavioral_face_mask	0.159	0.035	0.063	1.000	0.081	0.178	
behavioral_wash_hands	0.297	0.087	0.337	0.081	1.000	0.193	
behavioral_large_gatherings	0.258	-0.047	0.231	0.178	0.193	1.000	
behavioral_outside_home	0.247	-0.067	0.223	0.164	0.191	0.585	
behavioral_touch_face	0.249	0.084	0.333	0.104	0.364	0.254	
chronic_med_condition	0.096	-0.020	0.040	0.068	0.032	0.104	
child_under_6_months	0.048	0.023	-0.003	0.039	0.035	0.021	
health_worker	0.033	0.170	-0.001	0.069	0.053	-0.033	
opinion_seas_vacc_effective	0.235	0.081	0.116	0.044	0.139	0.080	
opinion_seas_risk	0.333	0.076	0.130	0.110	0.173	0.133	
opinion_seas_sick_from_vacc	0.223	-0.063	0.084	0.093	0.089	0.136	
household_adults	-0.019	0.018	0.015	0.013	0.004	-0.035	
household_children	0.050	0.048	0.038	0.003	0.043	-0.011	
seasonal_vaccine	0.160	0.121	0.079	0.051	0.114	0.065	
behav_score	0.394	0.040	0.616	0.336	0.579	0.704	
behav_to_risk	-0.027	-0.045	0.275	0.101	0.244	0.319	
older_65	0.018	-0.123	-0.020	0.002	-0.002	0.092	
high_risk_compl	0.090	-0.078	0.012	0.059	0.032	0.128	

```
In [62]: feats_to_drop = ['older_65', 'high_risk_compl']
df.drop(columns=feats_to_drop, axis=1, inplace=True)
df.head()
```

```
Out[62]:
```

	h1n1_concern	h1n1_knowledge	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavioral_large_gatherings	behavioral_outside_home	be
0	1.000	0.000	0.000	0.000	0.000	0.000	1.000	
1	3.000	2.000	1.000	0.000	1.000	0.000	1.000	
2	1.000	1.000	1.000	0.000	0.000	0.000	0.000	
3	1.000	1.000	1.000	0.000	1.000	1.000	0.000	
4	2.000	1.000	1.000	0.000	1.000	1.000	0.000	

We need to process numerical and categorical variables differently, but right now some categorical variables are still showing up as numeric because NaNs haven't been filled in with 'missing'. This can be done as part of the preprocessing pipeline.

```
In [63]: null_df = calculate_null_percentage(df)
miss_val_cols = list(null_df.loc[null_df['% Null']>0].index)
miss_val_cols
```

```
Out[63]: ['doctor_recc_seasonal',
'health_insurance',
'education',
'income_poverty',
'marital_status',
'rent_or_own',
'employment_status',
'employment_industry',
'employment_occupation']
```

these all need to have null values filled with 'missing' so they will all be changed to categorical features



```
In [64]: # list of all columns that are currently a object
obj_cols = list(df.select_dtypes('O').columns)
cat_cols = list(set(obj_cols + miss_val_cols))
cat_cols
```

```
Out[64]: ['census_msa',
'employment_industry',
'high_risk_cat',
'income_poverty',
'sex',
'health_insurance',
'employment_status',
'employment_occupation',
'education',
'doctor_recc_seasonal',
'age_group',
'marital_status',
'hhs_geo_region',
'race',
'rent_or_own']
```

```
In [65]: num_cols = [col for col in df.drop('seasonal_vaccine', axis=1).columns if col not in cat_cols]
num_cols
```

```
Out[65]: ['h1n1_concern',
'h1n1_knowledge',
'behavioral_avoidance',
'behavioral_face_mask',
'behavioral_wash_hands',
'behavioral_large_gatherings',
'behavioral_outside_home',
'behavioral_touch_face',
'chronic_med_condition',
'child_under_6_months',
'health_worker',
'opinion_seas_vacc_effective',
'opinion_seas_risk',
'opinion_seas_sick_from_vacc',
'household_adults',
'household_children',
'behav_score',
'behav_to_risk']
```

```
In [66]: # define target variable
target = 'seasonal_vaccine'

# separate of features (X) and target (y) for train-test-split
X = df.drop(columns=target, axis=1).copy()
y = df[target].copy()

# split the data into training and test sets prior to preprocessing
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)

## check for class imbalance across all sets of y
print(**original**, y.value_counts(normalize=True), '\n-----\n')
print(**y_train**, y_train.value_counts(normalize=True), '\n-----\n')
print(**y_test**, y_test.value_counts(normalize=True), '\n-----\n')

**original**
0    0.531
1    0.469
Name: seasonal_vaccine, dtype: float64
-----

**y_train**
0    0.534
1    0.466
Name: seasonal_vaccine, dtype: float64
-----

**y_test**
0    0.523
1    0.477
Name: seasonal_vaccine, dtype: float64
-----
```

LogisticRegression




```

In [67]: # transforming numerical columns
num_transformer = Pipeline(steps = [('scaler', StandardScaler())])

# transforming categorical columns and missing
cat_transformer = Pipeline(steps = [('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
                                   ('encoder', OneHotEncoder(handle_unknown='ignore',
                                                             sparse_output=False))])

preprocessing = ColumnTransformer(transformers=[('num', num_transformer, num_cols),
                                              ('cat', cat_transformer, cat_cols)])

modell = Pipeline([('preproc', preprocessing), ('model', LogisticRegression())])

modell.fit(X_train,y_train)

print(modell.score(X_test,y_test))
preds = modell.predict(X_test)
confusion_matrix(y_test,preds)

0.7825180433039294

/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessing.html)
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)
  n_iter_i = _check_optimize_result(

Out[67]: array([[2636,  628],
               [ 728, 2243]])

```



```
In [68]: param_grid = {
    'model__penalty': ['l1', 'l2'],
    'model__C': [0.1, 1.0, 10.0]
}

grid_search = GridSearchCV(model1, param_grid, cv=5)
grid_search.fit(X_train, y_train)

best_model = grid_search.best_estimator_

print("Best parameters:", grid_search.best_params_)
print("Best score:", grid_search.best_score_)

best_model_score = best_model.score(X_test, y_test)
print("Best model score:", best_model_score)

best_preds = best_model.predict(X_test)
confusion_matrix(y_test, best_preds)
```



```

/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```

n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html> (<https://scikit-learn.org/stable/modules/preprocessing.html>)

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)



```

dules/linear_model.html#logistic-regression)
    n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessin
g.html)
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/mo
dules/linear_model.html#logistic-regression)
    n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessin
g.html)
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/mo
dules/linear_model.html#logistic-regression)
    n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessin
g.html)
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/mo
dules/linear_model.html#logistic-regression)
    n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessin
g.html)
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/mo
dules/linear_model.html#logistic-regression)
    n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessin
g.html)
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/mo
dules/linear_model.html#logistic-regression)
    n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessin
g.html)
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/mo
dules/linear_model.html#logistic-regression)
    n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessin
g.html)
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression (https://scikit-learn.org/stable/mo
dules/linear_model.html#logistic-regression)
    n_iter_i = _check_optimize_result(
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/model_selection/_validation.py:378: FitFailedWarning:
15 fits failed out of a total of 30.
The score on these train-test partitions for these parameters will be set to nan.
If these failures are not expected, you can try to debug them by setting error_score='raise'.

Below are more details about the failures:
-----

```



```

15 fits failed with the following error:
Traceback (most recent call last):
  File "/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/model_selection/_validation.py", line 686, in
_fit_and_score
    estimator.fit(X_train, y_train, **fit_params)
  File "/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/pipeline.py", line 405, in fit
    self._final_estimator.fit(Xt, y, **fit_params_last_step)
  File "/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py", line 1162, in fit
    solver = _check_solver(self.solver, self.penalty, self.dual)
  File "/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py", line 54, in _check
_solver
    raise ValueError(
ValueError: Solver lbfgs supports only 'l2' or 'none' penalties, got l1 penalty.

    warnings.warn(some_fits_failed_message, FitFailedWarning)
/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/model_selection/_search.py:952: UserWarning: One or mo
re of the test scores are non-finite: [          nan 0.78090207          nan 0.78116945          nan 0.78090216]
    warnings.warn(

Best parameters: {'model__C': 1.0, 'model__penalty': 'l2'}
Best score: 0.7811694519609986
Best model score: 0.7825180433039294

/Users/jdapeman/anaconda3/lib/python3.10/site-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbf
gs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html (https://scikit-learn.org/stable/modules/preprocessin
g.html)
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression (https://scikit-learn.org/stable/mo
dules/linear\_model.html#logistic-regression)
    n_iter_i = _check_optimize_result(
Out[68]: array([[2636,  628],
               [ 728, 2243]])

```



```

In [69]: # Compute predicted probabilities for positive class
from sklearn.metrics import roc_auc_score
from sklearn.metrics import roc_curve


probs = best_model.predict_proba(X_test)[: , 1]

# Compute ROC curve values
fpr, tpr, thresholds = roc_curve(y_test, probs)

# Compute AUC score
auc_score = roc_auc_score(y_test, probs)

# Plot ROC curve
plt.plot(fpr, tpr, label='ROC curve (AUC = {:.2f})'.format(auc_score))
plt.plot([0, 1], [0, 1], 'k--') # Diagonal line
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc='lower right')
plt.show()

```



RandomForestClassifier

```

In [70]: # transforming numerical columns
num_transformer = Pipeline(steps = [('scaler', StandardScaler())])

# transforming categorical columns and missing
cat_transformer = Pipeline(steps = [('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
                                   ('encoder', OneHotEncoder(handle_unknown='ignore',
                                                             sparse_output=False))])

preprocessing = ColumnTransformer(transformers=[('num', num_transformer, num_cols),
                                              ('cat', cat_transformer, cat_cols)])

model12 = Pipeline([('preproc', preprocessing), ('model', RandomForestClassifier())])

model12.fit(X_train,y_train)

print(model12.score(X_test,y_test))
preds = model12.predict(X_test)
confusion_matrix(y_test,preds)

0.7797914995990377

```

```

Out[70]: array([[2662,  602],
               [ 771, 2200]])

```



```
In [71]: param_grid = {
        'model__criterion': ['gini', 'entropy'],
        'model__max_depth': [1, 2, 5, 10],
        'model__min_samples_split': [1, 5, 10, 20]
    }

    gs_RFC = GridSearchCV(model12, param_grid, cv=3)
    gs_RFC.fit(X_train,y_train)

    gs_RFC.best_params_
```

```
Out[71]: {'model__criterion': 'entropy',
          'model__max_depth': 10,
          'model__min_samples_split': 5}
```

KNeighborsClassifier

```
In [72]: # transforming numerical columns
num_transformer = Pipeline(steps = [('scaler', StandardScaler())])

# transforming categorical columns and missing
cat_transformer = Pipeline(steps = [('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
                                   ('encoder', OneHotEncoder(handle_unknown='ignore',
                                                             sparse_output=False))])

preprocessing = ColumnTransformer(transformers=[('num', num_transformer, num_cols),
                                              ('cat', cat_transformer, cat_cols)])

model13 = Pipeline([('preproc', preprocessing), ('model', KNeighborsClassifier())])

model13.fit(X_train,y_train)

print(model1.score(X_test,y_test))
preds = model1.predict(X_test)
confusion_matrix(y_test,preds)

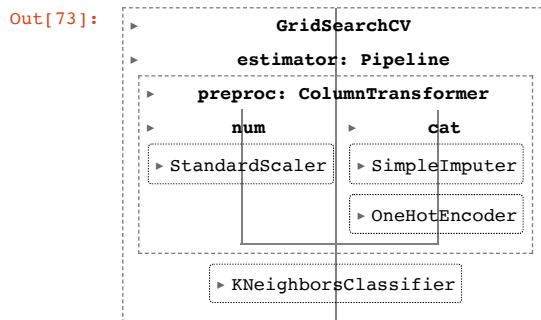
0.7825180433039294
```

```
Out[72]: array([[2636, 628],
               [ 728, 2243]])
```

```
In [73]: # fitting the model for grid search

param_grid = {
    'model__n_neighbors': range(2,10,2)}

gs_knn = GridSearchCV(model13, param_grid, cv=3)
gs_knn.fit(X_train,y_train)
```



```
In [74]: print(gs_knn.best_params_)
gs_knn.best_score_
```

```
{'model__n_neighbors': 8}
```

```
Out[74]: 0.7323563505680664
```




```
In [75]: # plot a ROC curve for the KNeighborsClassifier

# Get predicted probabilities for the positive class
probs = gs_knn.predict_proba(X_test)[: , 1]

# Calculate the false positive rate, true positive rate, and thresholds
fpr, tpr, thresholds = roc_curve(y_test, probs)

# Plot the ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, label='ROC curve')
plt.plot([0, 1], [0, 1], 'r--', label='Random Guess')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC)')
plt.legend()
plt.show()
```



DecisionTreeClassifier




```

In [76]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, roc_curve, auc
from sklearn.preprocessing import OneHotEncoder
from sklearn import tree

# transforming numerical columns
num_transformer = Pipeline(steps = [('scaler', StandardScaler())])

# transforming categorical columns and missing
cat_transformer = Pipeline(steps = [('imputer', SimpleImputer(strategy='constant', fill_value='missing')),
                                   ('encoder', OneHotEncoder(handle_unknown='ignore',
                                                             sparse_output=False))])

preprocessing = ColumnTransformer(transformers=[('num', num_transformer, num_cols),
                                              ('cat', cat_transformer, cat_cols)])

modell14 = Pipeline([('preproc', preprocessing), ('model', DecisionTreeClassifier())])

modell14.fit(X_train,y_train)

print(modell1.score(X_test,y_test))
preds = modell1.predict(X_test)
confusion_matrix(y_test,preds)

```

0.7825180433039294

```

Out[76]: array([[2636,  628],
               [ 728, 2243]])

```

```

In [77]: param_grid = {
    'model__criterion': ['gini', 'entropy'],
    'model__max_depth': [1, 2, 5, 10],
    'model__min_samples_split': [1, 5, 10, 20]
}

gs_tree = GridSearchCV(modell14, param_grid, cv=3)
gs_tree.fit(X_train,y_train)

gs_tree.best_params_

```

```


Out[77]: {'model__criterion': 'gini',
          'model__max_depth': 5,
          'model__min_samples_split': 20}

```



```
In [78]: # plot a ROC curve for the Decision tree classifier
fpr, tpr, thresholds = roc_curve(y_test, preds)
roc_auc = auc(fpr, tpr)

plt.figure()
plt.plot(fpr, tpr, label='ROC curve (area = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], 'k--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()
```



```
In [79]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

%matplotlib inline
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, AdaBoostClassifier
from sklearn.metrics import accuracy_score
```

Feature Importance

best model: LogisticRegression



```
In [80]: feature_importance = model1.named_steps['model'].coef_[0]
feature_names = model1.named_steps['preproc'].transformers_[1][1].named_steps['encoder'].get_feature_names_out(cat_cols)

# Create a DataFrame to store the feature importance
importance_df = pd.DataFrame({'Feature': np.concatenate((num_cols, feature_names)), 'Importance': feature_importance})

# Sort the DataFrame by importance values in descending order
importance_df = importance_df.sort_values('Importance', ascending=False)

# Select the top 5 features
top_features = importance_df.head(15)

# Print the top 5 features and their importance values
print(top_features)
```

	Feature	Importance
26	employment_industry_haxffmxo	1.530
63	employment_occupation_dcjcpih	1.530
91	doctor_recc_seasonal_1	0.894
97	age_group_65+ Years	0.841
12	opinion_seas_risk	0.711
31	employment_industry_msuufmds	0.692
11	opinion_seas_vacc_effective	0.646
84	employment_occupation_xzmlyyvj	0.411
54	health_insurance_1	0.362
34	employment_industry_phxvnwax	0.352
25	employment_industry_fcxhlnwr	0.237
29	employment_industry_mfikgejo	0.206
96	age_group_55 - 64 Years	0.200
105	hhs_geo_region_kbazzjca	0.191
21	employment_industry_arjwrjbj	0.190

```
In [81]: # Print top_features on a bar graph and selecting the relevant features
plt.figure(figsize=(10, 6))
sns.barplot(x='Importance', y='Feature', data=top_features)
plt.title('Top 5 Features')
plt.xlabel('Importance')
plt.ylabel('Feature')
plt.show()
```

